

```
/**

@mainpage 15-410 Project 3

@author Name1 (id1)
@author Name2 (id2)

Currently using README.dox as a bug tracker and TODO list.

BUGS:
- Currently accessing invalid memory with eip. This happens during call
  to smemalign (it seems). Eip is set to a value just above USER_MEM_START,
  and at that time no user memory has been allocated. This probably means some
interrupt
  is triggering a handler which is supposed to be located at that eip.

  Very puzzling. Maybe something to do with the console/keyboard interrupts????

CK1:
- Write gettid handler           [essential]
- Store tid at top of kernel stack [nice to have]
- Pass gettid_test1
- Cleanup unnecessary files and organize into folders [nice to have]

CK2:
- Synchronization design
- Context switch
- Proper physical memory management
- Trace facility                [nice to have]
- Fork or exec?

*/
```

```
/** @file asm_interrupt_handler.h
 *
 * @brief helper functions to save register values before calling C
 *        interrupt handler functions, and then restore registers and return
 *        to normal execution.
 */

#ifdef _P1_ASM_INTERRUPT_HANDLER_H_
#define _P1_ASM_INTERRUPT_HANDLER_H_

#include "../timer_driver.h"
#include "../keybd_driver.h"

/** @brief Saves all register values, calls timer_int_handler(), restores
 *        all register values and returns to normal execution.
 *
 * pusha is used to save all registers. Though not needed, %esp is also pushed
 * onto the stack for convenience of implementation. A call is made to
 * timer_int_handler(). When timer_int_handler() returns, popa is used
 * to restore all registers in the correct order, and iret to return to
 * normal execution prior to the interrupt.
 *
 * @return Void.
 */
void call_timer_int_handler(void);

/** @brief Saves all register values, calls keybd_int_handler(), restores
 *        all register values and returns to normal execution.
 *
 * pusha is used to save all registers. Though not needed, %esp is also pushed
 * onto the stack for convenience of implementation. A call is made to
 * keybd_int_handler(). When keybd_int_handler() returns, popa is used
 * to restore all registers in the correct order, and iret to return to
 * normal execution prior to the interrupt.
 *
 * @return Void.
 */
void call_keybd_int_handler(void);

#endif
```

```
/** @file asm_interrupt_handler.S
 * @brief Implementations for assembly functions
 * @author Nicklaus Choo (nchoo)
 * @bugs No known bugs
 */

.globl call_timer_int_handler

call_timer_int_handler:
    pusha /* Pushes all registers onto the stack */
    call timer_int_handler /* calls timer interrupt handler */
    popa /* Restores all registers onto the stack */
    iret /* Return to procedure before interrupt */

.globl call_keybd_int_handler

call_keybd_int_handler:
    pusha /* Pushes all registers onto the stack */
    call keybd_int_handler /* calls timer interrupt handler */
    popa /* Restores all registers onto the stack */
    iret /* Return to procedure before interrupt */
```

./kern/console.c

```

/** @file console.c
 * @brief Implements functions in console.h
 *
 * Since console.h contains comments, added comments will be written
 * with this format with preceding and succeeding --
 *
 * --
 * < additional comments >
 * --
 *
 * @author Andre Nascimento (anascime)
 * @author Nicklaus Choo (nchoo)
 */

#include <console.h> /* All console function prototypes */
#include <video_defines.h> /* CONSOLE_HEIGHT, CONSOLE_WIDTH */
#include <string.h> /* memmove () */
#include <assert.h> /* assert(), affirm() */
#include <x86/asm.h> /* outb() */

/* Default global console color. */
static int console_color = BGND_BLACK | FGND_WHITE;

/* Logical cursor row starts off at 0 */
static int cursor_row = 0;

/* Logical cursor col starts off at 0 */
static int cursor_col = 0;

/* Boolean for if cursor is hidden */
static int cursor_hidden = 0;

/* Background color mask to extract invalid set bits in color. FFFF FF00 */
#define INVALID_COLOR 0xFFFFF00

/** @brief Helper function to check if (row, col) is onscreen
 *
 * @param row Row index
 * @param col Col index
 * @return 1 if onscreen, 0 otherwise
 */
static int
onscreen(int row, int col) {
    return 0 <= row && row < CONSOLE_HEIGHT && 0 <= col && col < CONSOLE_WIDTH;
}

/** @brief Sets the hardware cursor to any row or column. Should only be called
 * by hide_cursor(), unhide_cursor(), set_cursor().
 *
 * Invariant for row and col is such that the cursor is only ever set to
 * be onscreen, or offscreen specifically at (CONSOLE_HEIGHT, CONSOLE_WIDTH).
 *
 * @param row Row to set hardware cursor to.
 * @param col Column to set hardware cursor to.
 * @return Void.
 */
static void
set_hardware_cursor( int row, int col )
{
    /* Only values for row, col if called by hide, unhide, set cursor functions */
    assert(onscreen(row, col) || (row == CONSOLE_HEIGHT && col == CONSOLE_WIDTH));

    /* Calculate offset in row major form */
    short hardware_cursor_offset = row * CONSOLE_WIDTH + col;

    /* Set lower 8 bits */
    outb(CRTC_IDX_REG, CRTC_CURSOR_LSB_IDX);

    outb(CRTC_DATA_REG, hardware_cursor_offset);

    /* Set upper 8 bits */
    outb(CRTC_IDX_REG, CRTC_CURSOR_MSB_IDX);
    outb(CRTC_DATA_REG, hardware_cursor_offset >> 8);
}

/** @brief Scrolls the terminal up by 1 line and ensures that the position of
 * the logical cursor remains fixed with respect to console output
 *
 * @return Void.
 */
static void
scroll( void )
{
    /* Move screen contents all up 1 row. */
    memmove((void *) CONSOLE_MEM_BASE,
            (void *) (CONSOLE_MEM_BASE + 2 * CONSOLE_WIDTH),
            2 * CONSOLE_WIDTH * (CONSOLE_HEIGHT - 1));

    /* The new last row should be an empty row of spaces */
    for (int col = 0; col < CONSOLE_WIDTH; col++) {
        draw_char(CONSOLE_HEIGHT - 1, col, ' ', console_color);
    }
    /* Cursor is stationary relative to output. */
    set_cursor(cursor_row - 1, cursor_col);
}

/** @brief Modification of the putbyte() specification which takes in
 * arguments for starting row and column. Useful for readline()
 *
 * putbyte() is a wrapper around this function.
 *
 * @param ch The character to print
 * @param start_rowp Pointer to starting row
 * @param start_colp Pointer to starting column
 * @return The input character
 */
int
scrolled_putbyte( char ch, int *start_rowp, int *start_colp )
{
    assert(start_rowp);
    assert(start_colp);
    assert(onscreen(*start_rowp, *start_colp));
    assert(onscreen(cursor_row, cursor_col));

    switch (ch) {
        case '\n': {
            /* Scroll if at screen bottom */
            if (cursor_row + 1 >= CONSOLE_HEIGHT) {
                scroll();
                // TODO you can scroll offscreen though
                *start_rowp -= 1;
            }
            /* Always update the cursor position relative to content */
            draw_char(cursor_row + 1, 0, ' ', console_color);
            set_cursor(cursor_row + 1, 0);
            break;
        }
        case '\r': {
            set_cursor(*start_rowp, *start_colp);
            break;
        }
        case '\b': {
            /* Not at leftmost column */

```

```

    if (cursor_col > 0) {

        /* Always draw space before moving else cursor blotted out */
        draw_char(cursor_row, cursor_col - 1, ' ', console_color);
        set_cursor(cursor_row, cursor_col - 1);

        /* At leftmost column, backspace goes to previous row if not top */
    } else {
        if (cursor_row > 0) {
            draw_char(cursor_row - 1, CONSOLE_WIDTH - 1, ' ',
                      console_color);
            set_cursor(cursor_row - 1, CONSOLE_WIDTH - 1);
        }
        break;
    }
}
default: {

    /* Print the character */
    draw_char(cursor_row, cursor_col, ch, console_color);

    /* If we are at the end of a line, set cursor on new line */
    if (cursor_col + 1 >= CONSOLE_WIDTH) {

        /* Scroll if necessary */
        if (cursor_row + 1 >= CONSOLE_HEIGHT) {
            scroll();
            *start_rowp -= 1;
        }
        /* Start printing below, update color if needed */
        char next_ch = get_char(cursor_row + 1, 0);
        draw_char(cursor_row + 1, 0, next_ch, console_color);
        set_cursor(cursor_row + 1, 0);
    } else {
        char next_ch = get_char(cursor_row, cursor_col + 1);
        draw_char(cursor_row, cursor_col + 1, next_ch, console_color);
        set_cursor(cursor_row, cursor_col + 1);
    }
    break;
}
}
assert(onscreen(cursor_row, cursor_col));
return ch;
}

/** @brief Prints character ch at the current location
 *   of the cursor.
 *
 *   If the character is a newline ('\n'), the cursor is moved
 *   to the beginning of the next line (scrolling if necessary).
 *   If the character is a carriage return ('\r'), the cursor is
 *   immediately reset to the beginning of the current line,
 *   causing any future output to overwrite any existing output
 *   on the line. If backspace ('\b') is encountered, the previous
 *   character is erased. See the main console.c description found
 *   on the handout web page for more backspace behavior.
 *   --
 *   We move the cursors as we write. Since there is no notion of a console
 *   prompt for putbyte(), the call to scrolled_putbyte() will have
 *   start_row == cursor_row and start_col == 0
 *   --
 *
 *   @param ch the character to print
 *   @return The input character
 */
int putbyte(char ch) {

    /* Get starting row and column, but set starting column to 0 */
    int start_row;
    int start_col;
    get_cursor(&start_row, &start_col);
    start_col = 0;

    return scrolled_putbyte(ch, &start_row, &start_col);
}

/** @brief Prints the string s, starting at the current
 *   location of the cursor.
 *
 *   If the string is longer than the current line, the
 *   string fills up the current line and then
 *   continues on the next line. If the string exceeds
 *   available space on the entire console, the screen
 *   scrolls up one line, and then the string
 *   continues on the new line. If '\n', '\r', and '\b' are
 *   encountered within the string, they are handled
 *   as per putbyte. If len is not a positive integer or s
 *   is null, the function has no effect.
 *
 *   @param s The string to be printed.
 *   @param len The length of the string s.
 *   @return Void.
 */
void
putbytes(const char *s, int len)
{
    affirm(s);
    if (len < 0) {
        return;
    }
    /* s is null string doesn't mean s == NULL */
    if (s[0] == '\0') {
        return;
    }
    for (int i = 0; i < len; i++) {
        char ch = s[i];
        putbyte(ch);
    }
}

/** @brief Changes the foreground and background color
 *   of future characters printed on the console.
 *
 *   If the color code is invalid, the function has no effect.
 *
 *   @param color The new color code.
 *   @return 0 on success or integer error code less than 0 if
 *   color code is invalid.
 */
int
set_term_color(int color)
{
    /* No effect if invalid color passed */
    if (color & INVALID_COLOR) {
        return -1;
    }
    /* Else set console_color */
    console_color = color;
    return 0;
}

/** @brief Writes the current foreground and background
 *   color of characters printed on the console
 *   into the argument color.
 *
 *   @param color The address to which the current color
 *   information will be written.

```

```

* @return Void.
*/
void
get_term_color( int* color )
{
    affirm(color);
    *color = console_color;
}

/** @brief Sets the position of the cursor to the
 *    position (row, col).
 *
 * Subsequent calls to putbytes should cause the console
 * output to begin at the new position. If the cursor is
 * currently hidden, a call to set_cursor() does not show
 * the cursor.
 * --
 * If cursor_hidden, the logical cursor is set without setting the
 * hardware cursor. This is because the hardware cursor is always
 * the one that is visible.
 *
 * Else, if there is a change in row, col, the logical cursor and the hardware
 * cursor are set.
 * --
 *
 * @param row The new row for the cursor.
 * @param col The new column for the cursor.
 * @return 0 on success or integer error code less than 0 if
 *         cursor location is invalid.
 */
int
set_cursor( int row, int col )
{
    assert(onscreen(cursor_row, cursor_col));
    /* set logical cursor */
    if (onscreen(row, col)) {
        cursor_row = row;
        cursor_col = col;

        /* If cursor not hidden and change in row, col, set hardware cursor */
        if (!cursor_hidden && (cursor_row != row || cursor_col != col)) {
            set_hardware_cursor(row, col);
        }
        assert(onscreen(cursor_row, cursor_col));
        return 0;
    }
    /* cursor location is invalid, do nothing and return -1 */
    assert(onscreen(cursor_row, cursor_col));
    return -1;
}

/** @brief Writes the current position of the cursor
 *    into the arguments row and col.
 *
 * --
 * Only writes to row, col if they are non-null, throws affirm() error
 * otherwise.
 * --
 *
 * @param row The address to which the current cursor
 *            row will be written.
 * @param col The address to which the current cursor
 *            column will be written.
 * @return Void.
 */
void
get_cursor( int* row, int* col )

```

```

{
    affirm(row);
    affirm(col);
    *row = cursor_row;
    *col = cursor_col;
}

/** @brief Shows the cursor.
 *
 * If the cursor is already shown, the function has no effect.
 * --
 * Hides the cursor by setting the hardware cursor to
 * (CONSOLE_HEIGHT, CONSOLE_WIDTH) and toggles cursor_hidden to true i.e. 1.
 * Note that this function is idempotent.
 * --
 *
 * @return Void.
 */
void
hide_cursor( void )
{
    assert(onscreen(cursor_row, cursor_col));
    set_hardware_cursor(CONSOLE_HEIGHT, CONSOLE_WIDTH);
    cursor_hidden = 1;
}

/** @brief Shows the cursor.
 *
 * If the cursor is already shown, the function has no effect.
 * --
 * Shows the cursor by setting the hardware cursor to the current location
 * of the logical cursor and toggles cursor_hidden to false i.e. 0.
 * Note that this function is idempotent.
 * --
 *
 * @return Void.
 */
void
show_cursor( void )
{
    assert(onscreen(cursor_row, cursor_col));
    set_hardware_cursor(cursor_row, cursor_col);
    cursor_hidden = 0;
}

/** @brief Clears the entire console.
 *
 * The cursor is reset to the first row and column
 *
 * @return Void.
 */
void
clear_console( void )
{
    /* Replace everything onscreen with a blank space */
    for (size_t row = 0; row < CONSOLE_HEIGHT; row++) {
        for (size_t col = 0; col < CONSOLE_WIDTH; col++) {
            draw_char(row, col, ' ', console_color);
        }
    }
    /* Set cursor to the top left corner */
    set_cursor(0, 0);
}

/** @brief Prints character ch with the specified color
 *
 * at position (row, col).

```

```
*
* If any argument is invalid, the function has no effect.
*
* @param row The row in which to display the character.
* @param col The column in which to display the character.
* @param ch The character to display.
* @param color The color to use to display the character.
* @return Void.
*/
void
draw_char( int row, int col, int ch, int color )
{
    /* If row or col out of range, invalid row, no effect. */
    if (!onscreen(row, col)) {
        return;
    }
    /* If background color not supported, invalid color, no effect. */
    if (color & INVALID_COLOR) {
        return;
    }
    /* All arguments valid, draw character with specified color */
    char *chp = (char *) (CONSOLE_MEM_BASE + 2*(row * CONSOLE_WIDTH + col));
    *chp = ch;
    *(chp + 1) = color;
}

/** @brief Returns the character displayed at position (row, col).
*
* --
* Offscreen characters are always NULL or simply 0 by default.
* --
*
* @param row Row of the character.
* @param col Column of the character.
* @return The character at (row, col).
*/
char
get_char( int row, int col )
{
    /* If out of range, return 0. */
    if (!onscreen(row, col)) {
        return 0;
    }
    /* Else return char at row, col. */
    return *(char *) (CONSOLE_MEM_BASE + 2*(row * CONSOLE_WIDTH + col));
}
```

./kern/console_driver.c

```

/** @file console_driver.c
 * @brief Contains console driver functions implementint the plkern interface,
 * and other helper functions not in the interface as well.
 *
 * For interface functions as declared in plkern.h, additional comment blocks
 * are preceeded and succeeded by '--'.
 *
 * @author Nicklaus Choo (nchoo)
 * @bug No known bugs
 */

#include <limits.h> /* INT_MAX */
#include <ctype.h> /* isprint() */
#include <plkern.h> /* client interface functions */
#include <stddef.h> /* NULL */
#include <assert.h> /* assert() */
#include <x86/asm.h> /* outb() */
#include <string.h> /* memmove () */

/* Default global console color. */
static int console_color = BGND_BLACK | FGND_WHITE;

/* Logical cursor row starts off at 0 */
static int cursor_row = 0;

/* Logical cursor col starts off at 0 */
static int cursor_col = 0;

/* Boolean for if cursor is hidden */
static int cursor_hidden = 0;

/* Background color mask to extract invalid set bits in color. FFFF FF00 */
#define INVALID_COLOR 0xFFFFF00

/*****
 *
 * Internal helper functions
 *
 *****/

/** @brief Checks if row and col are within console screen dimensions.
 *
 * @param row Row to check
 * @param col Col to check
 * @return value greater than 1 if true, 0 otherwise
 */
int
onscreen(int row, int col) {
    return 0 <= row && row < CONSOLE_HEIGHT && 0 <= col && col < CONSOLE_WIDTH;
}

/** @brief Sets the hardware cursor to any row or column. Should only be called
 * by hide_cursor(), unhide_cursor(), set_cursor().
 *
 * Invariant for row and col is such that the cursor is only ever set to
 * be onscreen, or offscreen specifically at (CONSOLE_HEIGHT, CONSOLE_WIDTH).
 *
 * @param row Row to set hardware cursor to.
 * @param col Column to set hardware cursor to.
 * @return Void.
 */
void
set_hardware_cursor(int row, int col) {

    /* Only values for row, col if called by hide, unhide, set cursor functions */
    assert(onscreen(row, col) || (row == CONSOLE_HEIGHT && col == CONSOLE_WIDTH));

    /* Calculate offset in row major form */
    short hardware_cursor_offset = row * CONSOLE_WIDTH + col;

    /* Set lower 8 bits */
    outb(CRTC_IDX_REG, CRTC_CURSOR_LSB_IDX);
    outb(CRTC_DATA_REG, hardware_cursor_offset);

    /* Set upper 8 bits */
    outb(CRTC_IDX_REG, CRTC_CURSOR_MSB_IDX);
    outb(CRTC_DATA_REG, hardware_cursor_offset >> 8);

    return;
}

/** @brief Scrolls the terminal up by 1 line and ensures that the position of
 * the logical cursor remains fixed with respect to console output
 *
 * @return Void.
 */
void
scroll(void) {

    /* Move screen contents all up 1 row. */
    memmove((void *) CONSOLE_MEM_BASE,
            (void *) (CONSOLE_MEM_BASE + 2 * CONSOLE_WIDTH),
            2 * CONSOLE_WIDTH * (CONSOLE_HEIGHT - 1));

    /* The new last row should be an empty row of spaces */
    for (size_t col = 0; col < CONSOLE_WIDTH; col++) {
        draw_char(CONSOLE_HEIGHT - 1, col, ' ', console_color);
    }

    /* Cursor is stationary relative to output. */
    set_cursor(cursor_row - 1, cursor_col);
}

/** @brief Modification of the putbyte() specification to record if there was
 * a screen scroll of 1 line after we put a byte on the screen.
 *
 * putbyte() is a wrapper around this function. Useful for readline() in
 * the keyboard driver.
 *
 * @param ch The character to print
 * @param scrolled Must be 0. Set to 1 if scrolled a line, unchanged otherwise
 * @return The input character
 */
int
_putbyte(char ch, int *scrolled) {
    assert(scrolled != NULL);
    assert(*scrolled == 0);
    assert(onscreen(cursor_row, cursor_col));

    /* Get current cursor position */
    char ogch = ch;
    int row, col;
    get_cursor(&row, &col);

    switch (ch) {
        case '\n': {

            /* Scroll if at screen bottom */
            if (cursor_row + 1 >= CONSOLE_HEIGHT) {
                scroll();
                *scrolled = 1;
            }

            /* Always update the cursor position relative to content */
            set_cursor(cursor_row + 1, 0);
            break;
        }
    }
}

```


./kern/console_driver.c

```

}
case '\r': {
    set_cursor(cursor_row, 0);
    break;
}
case '\b': {

    /* Not at leftmost column */
    if (cursor_col != 0) {

        /* Always draw space before moving else cursor will be blotted out */
        draw_char(cursor_row, cursor_col - 1, ' ', console_color);
        set_cursor(cursor_row, cursor_col - 1);

        /* At leftmost column, backspace goes to previous row */
    } else {
        draw_char(cursor_row - 1, CONSOLE_WIDTH - 1, ' ', console_color);
        set_cursor(cursor_row - 1, CONSOLE_WIDTH - 1);
    }
    break;
}
default: {

    /* Print the character, unprintable characters return -1 */
    if (!isprint(ch)) return -1;
    draw_char(cursor_row, cursor_col, ch, console_color);

    /* If we are at the end of a line, set cursor on new line */
    if (cursor_col + 1 >= CONSOLE_WIDTH) {

        /* Scroll if necessary */
        if (cursor_row + 1 >= CONSOLE_HEIGHT) {
            scroll();
            *scrolled = 1;
        }
        /* Start printing below */
        set_cursor(cursor_row + 1, 0);
    } else {
        set_cursor(cursor_row, cursor_col + 1);
    }
    break;
}
}
assert(onscreen(row, col));
return ogch;
}

/*****
/* Console interface functions
*****/

/** @brief Prints character ch at the current location
 * of the cursor.
 *
 * If the character is a newline ('\n'), the cursor is moved
 * to the beginning of the next line (scrolling if necessary).
 * If the character is a carriage return ('\r'), the cursor is
 * immediately reset to the beginning of the current line,
 * causing any future output to overwrite any existing output
 * on the line. If backspace ('\b') is encountered, the previous
 * character is erased. See the main console.c description found
 * on the handout web page for more backspace behavior.
 * --
 * Unprintable characters return -1.
 * We move the cursors as we write.

```

```

 * --
 * @param ch the character to print
 * @return The input character
 */
int putbyte(char ch) {

    /* Wrapper for _putbyte(). We don't care in putbyte() if we scrolled */
    int scrolled = 0;
    return _putbyte(ch, &scrolled);
}

/** @brief Prints the string s, starting at the current
 * location of the cursor.
 *
 * If the string is longer than the current line, the
 * string fills up the current line and then
 * continues on the next line. If the string exceeds
 * available space on the entire console, the screen
 * scrolls up one line, and then the string
 * continues on the new line. If '\n', '\r', and '\b' are
 * encountered within the string, they are handled
 * as per putbyte. If len is not a positive integer or s
 * is null, the function has no effect.
 *
 * @param s The string to be printed.
 * @param len The length of the string s.
 * @return Void.
 */
void
putbytes(const char *s, int len)
{
    if (len < 0 || s == NULL) return;
    for (size_t i = 0; i < len; i++) {
        char ch = s[i];
        putbyte(ch);
    }
}

/** @brief Prints character ch with the specified color
 * at position (row, col).
 *
 * If any argument is invalid, the function has no effect.
 * --
 * Note that we only draw printable characters. To draw something is to make
 * it to be seen. If a character cannot be seen/printed it cannot be
 * drawn and thus we do not draw unprintable characters
 * --
 * @param row The row in which to display the character.
 * @param col The column in which to display the character.
 * @param ch The character to display.
 * @param color The color to use to display the character.
 * @return Void.
 */
void draw_char(int row, int col, int ch, int color) {

    /* If row out of range, invalid row, no effect. */
    if (!(0 <= row && row < CONSOLE_HEIGHT)) return;

    /* If col out of range, invalid col, no effect. */
    if (!(0 <= col && col < CONSOLE_WIDTH)) return;

    /* If ch not printable, invalid ch, no effect. */
    if (!isprint(ch)) return;

    /* If background color not supported, invalid color, no effect. */
    if (color & INVALID_COLOR) return;

```

./kern/console_driver.c

```

/* All arguments valid, draw character */
*(char *) (CONSOLE_MEM_BASE + 2*(row * CONSOLE_WIDTH + col)) = ch;
*(char *) (CONSOLE_MEM_BASE + 2*(row * CONSOLE_WIDTH + col) + 1) = color;
}

/** @brief Returns the character displayed at position (row, col).
 * @param row Row of the character.
 * @param col Column of the character.
 * @return The character at (row, col).
 */
char get_char(int row, int col) {

    /* If out of range, return '\0'. */
    if (!(0 <= row && row < CONSOLE_HEIGHT)
        || !(0 <= col && col < CONSOLE_WIDTH)) {
        return '\0';
    }
    /* Else return char at row, col. */
    return *(char *) (CONSOLE_MEM_BASE + 2*(row * CONSOLE_WIDTH + col));
}

/** @brief Sets the position of the cursor to the
 * position (row, col).
 *
 * Subsequent calls to putbytes should cause the console
 * output to begin at the new position. If the cursor is
 * currently hidden, a call to set_cursor() does not show
 * the cursor.
 * --
 * If cursor_hidden, the logical cursor is set without setting the
 * hardware cursor. This is because the hardware cursor is always
 * the one that is visible.
 *
 * Else, the logical cursor and the hardware cursor are set.
 * --
 * @param row The new row for the cursor.
 * @param col The new column for the cursor.
 * @return 0 on success or integer error code less than 0 if
 * cursor location is invalid.
 */
int set_cursor(int row, int col) {

    /* set logical cursor */
    if (onscreen(row, col)) {
        cursor_row = row;
        cursor_col = col;

        /* If cursor is not hidden, set the hardware cursor */
        if (!cursor_hidden) set_hardware_cursor(row, col);
        return 0;
    }
    /* cursor location is invalid, do nothing and return -1 */
    return -1;
}

/** @brief Writes the current position of the cursor
 * into the arguments row and col.
 *
 * --
 * Only writes to row, col if they are non-null
 * --
 * @param row The address to which the current cursor
 * row will be written.
 * @param col The address to which the current cursor
 * column will be written.
 * @return Void.
 */

```

```

void get_cursor(int* row, int* col) {
    if (row != NULL) *row = cursor_row;
    if (col != NULL) *col = cursor_col;
    return;
}

/** @brief Shows the cursor.
 *
 * If the cursor is already shown, the function has no effect.
 * --
 * Hides the cursor by setting the hardware cursor to
 * (CONSOLE_HEIGHT, CONSOLE_WIDTH) and toggles cursor_hidden to true i.e. 1.
 * Note that this function is idempotent.
 * --
 * @return Void.
 */
void hide_cursor(void) {
    assert(onscreen(cursor_row, cursor_col));
    set_hardware_cursor(CONSOLE_HEIGHT, CONSOLE_WIDTH);
    cursor_hidden = 1;
    return;
}

/** @brief Shows the cursor.
 *
 * If the cursor is already shown, the function has no effect.
 * --
 * Shows the cursor by setting the hardware cursor to the current location
 * of the logical cursor and toggles cursor_hidden to false i.e. 0.
 * Note that this function is idempotent.
 * --
 * @return Void.
 */
void show_cursor(void) {
    assert(onscreen(cursor_row, cursor_col));
    set_hardware_cursor(cursor_row, cursor_col);
    cursor_hidden = 0;
    return;
}

/** @brief Writes the current foreground and background
 * color of characters printed on the console
 * into the argument color.
 * @param color The address to which the current color
 * information will be written.
 * @return Void.
 */
void get_term_color(int* color) {
    if (color != NULL) *color = console_color;
}

/** @brief Changes the foreground and background color
 * of future characters printed on the console.
 *
 * If the color code is invalid, the function has no effect.
 *
 * @param color The new color code.
 * @return 0 on success or integer error code less than 0 if
 * color code is invalid.
 */
int set_term_color(int color) {

    /* No effect if invalid color passed */
    if (color & INVALID_COLOR) return -1;

    /* Else set console_color */
    console_color = color;
}

```

```
    return 0;
}

/** @brief Clears the entire console.
 *
 * The cursor is reset to the first row and column
 *
 * @return Void.
 */
void clear_console(void) {
    for (size_t row = 0; row < CONSOLE_HEIGHT; row++) {
        for (size_t col = 0; col < CONSOLE_WIDTH; col++) {
            draw_char(row, col, ' ', console_color);
        }
    }
    /* Set cursor to the top left corner */
    set_cursor(0, 0);
}
```

./kern/game.c

```

/** @file game.c
 * @brief A kernel with timer, keyboard, console support which serves mainly
 * as a big file of test functions.
 *
 * This file contains the kernel's main() function.
 *
 * It sets up the drivers and starts the game. It contains test code to
 * exercise the kernel
 *
 * @author Nicklaus Choo (nchoo)
 * @bug No known bugs.
 */
#include <assert.h>
#include "../spec/plkern.h"

/* libc includes. */
#include <stdio.h>
#include <simics.h>          /* lprintf() */
#include <malloc.h>

/* multiboot header file */
#include <multiboot.h>      /* boot_info */

/* memory includes. */
#include <lmm.h>            /* lmm_remove_free() */

/* x86 specific includes */
#include <x86/seg.h>        /* install_user_segs() */
#include <x86/interrupt_defines.h> /* interrupt_setup() */
#include <x86/asm.h>        /* enable_interrupts() */

#include <string.h>

volatile static int __kernel_all_done = 0;

/* Think about where this declaration
 * should be... probably not here!
 */
void tick(unsigned int numTicks) {
    //lprintf("numTicks: %d\n", numTicks);
}

void test_scroll(void) {
    for (size_t i = 0; i < CONSOLE_HEIGHT + 1; i++) {
        printf("%d\n", i);
    }
}

void test_putbyte(void) {
    for (size_t i = 0; i < CONSOLE_HEIGHT; i++) {
        draw_char(i, 1, (i % 10) + '0', FGND_RED);
    }
    printf("a");
}

/** @brief Runs a few assert statements to test draw_char() and
 * get_char() functions
 *
 *
 */
void test_draw_char_get_char(void) {
    lprintf("Testing draw_char() and get_char()");
    draw_char(0,0, 'A', FGND_RED | BGND_BLACK);
    assert(get_char(0,0) == 'A');

    draw_char(0, CONSOLE_WIDTH - 1, 'B', FGND_WHITE | BGND_BLUE);
    assert(get_char(0, CONSOLE_WIDTH - 1) == 'B');

```

```

    draw_char(CONSOLE_HEIGHT - 1, CONSOLE_WIDTH - 1, 'C', FGND_RED | BGND_GREEN);
    assert(get_char(CONSOLE_HEIGHT - 1, CONSOLE_WIDTH - 1) == 'C');

    draw_char(CONSOLE_HEIGHT - 1, 0, 'D', FGND_YLLW | BGND_CYAN);
    assert(get_char(CONSOLE_HEIGHT - 1, 0) == 'D');

    /* Offscreen row drawing has no effect */
    draw_char(CONSOLE_HEIGHT, 0, 'E', FGND_YLLW | BGND_CYAN);
    assert(get_char(CONSOLE_HEIGHT, 0) == '\0');

    /* Offscreen col drawing has no effect */
    draw_char(CONSOLE_HEIGHT - 1, CONSOLE_WIDTH, 'E', FGND_YLLW | BGND_CYAN);
    assert(get_char(CONSOLE_HEIGHT - 1, CONSOLE_WIDTH) == '\0');

    /* Invalid background color drawing has no effect */
    draw_char(CONSOLE_HEIGHT - 1, 0, 'F', FGND_YLLW | 0x190);
    assert(get_char(CONSOLE_HEIGHT - 1, 0) == 'D');

    /* Invalid unprintable character drawing has becomes ? */
    draw_char(CONSOLE_HEIGHT - 1, 0, '\6', FGND_YLLW | BGND_RED);
    assert(get_char(CONSOLE_HEIGHT - 1, 0) == 'D');

    lprintf("Passed draw_char() and get_char()");
    return;
}

void test_cursor(void) {

    lprintf("Testing: set_cursor()");
    /* Out of bounds checks */
    assert(set_cursor(CONSOLE_HEIGHT, 0) == -1);
    assert(set_cursor(-1, 0) == -1);
    assert(set_cursor(0, CONSOLE_WIDTH) == -1);
    assert(set_cursor(0, -1) == -1);

    /* In bounds checks */
    assert(set_cursor(0, 0) == 0);
    assert(set_cursor(0, CONSOLE_WIDTH - 1) == 0);
    assert(set_cursor(CONSOLE_HEIGHT - 1, CONSOLE_WIDTH - 1)
           == 0);
    assert(set_cursor(CONSOLE_HEIGHT - 1, 0)
           == 0);
    assert(set_cursor(0, 0) == 0);
    lprintf("Passed: set_cursor()");
}

/** @brief Kernel entrypoint.
 *
 * This is the entrypoint for the kernel. It simply sets up the
 * drivers and passes control off to game_run().
 *
 * @return Does not return
 */
int kernel_main(mbinfo_t *mbinfo, int argc, char **argv, char **envp)
{
    /*
     * Initialize device-driver library.
     */
    int res = handler_install(tick);
    lprintf("res of handler_install: %d", res);

    /*
     * When kernel_main() begins, interrupts are DISABLED.
     * You should delete this comment, and enable them --
     * when you are ready.
     */

```

```
    */
    printf("h");

    lprintf( "Hello from a brand new kernel!" );

    char * badguy = (char *) 0xdeadd00d;
    char * s = "hello";
    lprintf("badguy: 0x%08lx, s: 0x%08lx", (long) badguy, (long) s);
    //putbytes((char *), 3);

    test_draw_char_get_char();
    //test_putbyte();

    test_cursor();

    //clear_console();

    //printf("Hello Mom!");
    //
    test_scroll();

    clear_console();

    putbytes("carriage return should bring the cursor to the front of this line.\r\n", 67);

    int tohide = 0;
    while (!__kernel_all_done) {
        if (tohide) hide_cursor();
        else show_cursor();
        int n = CONSOLE_HEIGHT * CONSOLE_WIDTH;
        char s[n];
        int res = readline(s, n);
        lprintf("characters read: %d, '%s'", res, s);
        putbytes("how many characters to read next", 33);
        res = readline(s, n);
        lprintf("characters read: %d, '%s'", res, s);

        tohide = !tohide;

        continue;
    }

    return 0;
}
```

```
/** @ file console_driver.h
 * @ brief Contains helpful internal functions for other drivers to use.
 *       Mainly the keyboard driver
 *
 * @author Nicklaus Choo (nchoo)
 * @bug No known bugs.
 */
#ifndef _P1_CONSOLE_DRIVER_H_
#define _P1_CONSOLE_DRIVER_H_

int _putbyte(char ch, int *scrolled);

#endif
```

./kern/inc/console.h

```

/*
 *
 * # #
 * ## # #### ##### # #### #####
 * # # # # # # # #
 * # # # # # # # #
 * # # # # # # # #
 * # ## # # # # #
 * # # #### # # #### #####
 *
 * Now that it's P3 instead of P1 you are allowed
 * to edit this file if it suits you.
 *
 * Please delete this notice.
 */

/** @file console.h
 * @brief Function prototypes for the console driver.
 *
 * This contains the prototypes and global variables for the console
 * driver
 *
 * @author Michael Berman (mberman)
 * @bug No known bugs.
 */

#ifndef _CONSOLE_H
#define _CONSOLE_H

#include <video_defines.h>

/** @brief Prints character ch at the current location
 * of the cursor.
 *
 * If the character is a newline ('\n'), the cursor is
 * be moved to the beginning of the next line (scrolling if necessary). If
 * the character is a carriage return ('\r'), the cursor
 * is immediately reset to the beginning of the current
 * line, causing any future output to overwrite any existing
 * output on the line. If backspace ('\b') is encountered,
 * the previous character is erased. See the main console.c description
 * for more backspace behavior.
 *
 * @param ch the character to print
 * @return The input character
 */
int putbyte( char ch );

/** @brief Prints the string s, starting at the current
 * location of the cursor.
 *
 * If the string is longer than the current line, the
 * string fills up the current line and then
 * continues on the next line. If the string exceeds
 * available space on the entire console, the screen
 * scrolls up one line, and then the string
 * continues on the new line. If '\n', '\r', and '\b' are
 * encountered within the string, they are handled
 * as per putbyte. If len is not a positive integer or s
 * is null, the function has no effect.
 *
 * @param s The string to be printed.
 * @param len The length of the string s.
 * @return Void.
 */
void putbytes(const char* s, int len);

/** @brief Changes the foreground and background color
 * of future characters printed on the console.
 *
 * If the color code is invalid, the function has no effect.
 *
 * @param color The new color code.
 * @return 0 on success or integer error code less than 0 if
 * color code is invalid.
 */
int set_term_color(int color);

/** @brief Writes the current foreground and background
 * color of characters printed on the console
 * into the argument color.
 *
 * @param color The address to which the current color
 * information will be written.
 * @return Void.
 */
void get_term_color(int* color);

/** @brief Sets the position of the cursor to the
 * position (row, col).
 *
 * Subsequent calls to putbytes should cause the console
 * output to begin at the new position. If the cursor is
 * currently hidden, a call to set_cursor() does not show
 * the cursor.
 *
 * @param row The new row for the cursor.
 * @param col The new column for the cursor.
 * @return 0 on success or integer error code less than 0 if
 * cursor location is invalid.
 */
int set_cursor(int row, int col);

/** @brief Writes the current position of the cursor
 * into the arguments row and col.
 *
 * @param row The address to which the current cursor
 * row will be written.
 * @param col The address to which the current cursor
 * column will be written.
 * @return Void.
 */
void get_cursor(int* row, int* col);

/** @brief Hides the cursor.
 *
 * Subsequent calls to putbytes do not cause the
 * cursor to show again.
 *
 * @return Void.
 */
void hide_cursor(void);

/** @brief Shows the cursor.
 *
 * If the cursor is already shown, the function has no effect.
 *
 * @return Void.
 */
void show_cursor(void);

/** @brief Clears the entire console.
 *
 * The cursor is reset to the first row and column
 */

```

```
* @return Void.
*/
void clear_console(void);

/** @brief Prints character ch with the specified color
 *   at position (row, col).
 *
 * If any argument is invalid, the function has no effect.
 *
 * @param row The row in which to display the character.
 * @param col The column in which to display the character.
 * @param ch The character to display.
 * @param color The color to use to display the character.
 * @return Void.
 */
void draw_char(int row, int col, int ch, int color);

/** @brief Returns the character displayed at position (row, col).
 *   @param row Row of the character.
 *   @param col Column of the character.
 *   @return The character at (row, col).
 */
char get_char(int row, int col);

/* helper for keybd */
int scrolled_putbyte( char ch, int *start_rowp, int *start_colp );

#endif /* _CONSOLE_H */
```


03/06/22
14:54:23

./kern/inc/install_handler.h

1

```
#ifndef _INSTALL_HANDLER_H_
#define _INSTALL_HANDLER_H_
int handler_install(void (*tickback)(unsigned int));
#endif
```

```
/** @file iret_travel.h
 *  @brief Asm facilities for moving to user mode */

#ifndef _IRET_TRAVEL_H
#define _IRET_TRAVEL_H

#include <stdint.h> /* uint32_t */

/** @brief Sets processor state. This can be used for "travelling"
 *      to user mode. Only to be called from kernel mode.
 *
 *      This function does not return
 */
void iret_travel( uint32_t ss, uint32_t esp,
                  uint32_t eflags, uint32_t cs, uint32_t eip );

#endif /* _IRET_TRAVEL_H */
```

```
/** @file keybd_driver.h
 * @brief Contains functions that can be called by interrupt handler assembly
 * wrappers
 *
 * @author Nicklaus Choo (nchoo)
 * @bugs No known bugs.
 */

# ifndef _P1_KEYBD_DRIVER_H_
# define _P1_KEYBD_DRIVER_H_

#include <stdint.h>

void init_keybd(void);
void keybd_int_handler(void);
int readline(char *buf, int len);

typedef int aug_char;
typedef uint8_t raw_byte;

/** @brief unbounded circular array heavily inspired by 15-122 unbounded array
 *
 * first is the index of the earliest unread element in the buffer. Once
 * all functions that will ever need to read a buffer element has read
 * the element at index 'first', first++ modulo limit.
 *
 * last is one index after the latest element added to the buffer. Whenever
 * an element is added to the buffer, last++ modulo limit.
 */
typedef struct {
    uint32_t size; /* 0 <= size && size < limit */
    uint32_t limit; /* 0 < limit */
    uint32_t first; /* if first <= last, then size == last - first */
    uint32_t last; /* else last < first, then size == limit - first + last */
    raw_byte *data;
} uba;

int is_uba(uba *arr);
uba *uba_new(int limit);
uba *uba_resize(uba *arr);
void uba_add(uba *arr, uint8_t elem);
uint8_t uba_rem(uba *arr);
int uba_empty(uba *arr);

#endif
```

```
/* The 15-410 kernel project
 *
 * loader.h
 *
 * Structure definitions, #defines, and function prototypes
 * for the user process loader.
 */

#ifndef _LOADER_H
#define _LOADER_H

/* --- Prototypes --- */

int getbytes( const char *filename, int offset, int size, char *buf );

int execute_user_program( const char *fname, int argc, char **argv );

#endif /* _LOADER_H */
```

```
/** @ file mem_manager.h
 * @ brief Contains helpful internal functions for other drivers to use.
 *        Mainly the keyboard driver
 *
 * @author Nicklaus Choo (nchoo)
 * @bug No known bugs.
 */
#ifndef _P1_MEM_MANAGER_H_
#define _P1_MEM_MANAGER_H_

void paging_on( void );
void paging_off( void );

int vm_init( void );

#endif
```

```
/** @brief Header file with VM API. */

#ifndef _MEMORY_MANAGER_H
#define _MEMORY_MANAGER_H

#include <elf_410.h> /* simple_elf_t */
#include <stdint.h> /* uint32_t */

int vm_init( void );
int vm_task_new ( void *ptd, simple_elf_t *elf,
                 uint32_t stack_lo, uint32_t stack_len );
void vm_enable_task( void *ptd );
void enable_write_protection( void );
void disable_write_protection( void );
int vm_new_pages ( void *ptd, void *base, int len );

#endif /* _MEMORY_MANAGER_H */
```

```
/** @brief Facilities for task creation and management.
 * Also provides context switching functionality. */

#ifndef _TASK_MANAGER_H
#define _TASK_MANAGER_H

#include <stdint.h> /* uint32_t */
#include <elf_410.h> /* simple_elf_t */

typedef struct pcb pcb_t;
typedef struct tcb tcb_t;

/* TODO: Do we even want PCB and TCB to be available to others? */
/** @brief Task control block */
struct pcb {
    void *ptd; // page table directory
    tcb_t *first_thread; // First thread in linked list

    int pid;

    int prepared; // Whether this task's VM has been initialized
};

/** @brief Thread control block */
struct tcb {
    pcb_t *owning_task;
    tcb_t *next_thread; // Embeded linked list of threads from same task

    int tid;

    /* Stack info. Needed for resuming execution.
     * General purpose registers, program counter
     * are stored on stack pointed to by esp. */
    int esp;
};

int task_new( int pid, int tid, simple_elf_t *elf );
int task_prepare( int pid );
void task_set( int tid, uint32_t esp, uint32_t entry_point );
void task_switch( int pid );

#endif /* _TASK_MANAGER_H */
```

```
/** @file timer_driver.h
 * @brief Contains internal functions that implement the timer driver.
 *
 * This header is to be included by install_handler.c
 *
 * @author Nicklaus Choo (nchoo)
 * @bug No known bugs.
 */

#ifndef _P1_TIMER_DRIVER_H_
#define _P1_TIMER_DRIVER_H_

void init_timer(void (*tickback)(unsigned int));

#endif
```


./kern/install_handler.c

```

/** @file install_handler.c
 * @brief Contains functions that install the timer and keyboard handlers
 *
 * @author Nicklaus Choo (nchoo)
 * @bug No known bugs.
 */

#include <x86/asm.h> /* idt_base() */
#include <assert.h> /* assert() */
#include <x86/interrupt_defines.h>
#include <x86/timer_defines.h> /* TIMER_IDT_ENTRY */
#include <x86/seg.h> /* SEGSEL_KERNEL_CS */
#include <x86/keyhelp.h>
#include <stddef.h> /* NULL */
#include <simics.h> /* lprintf */
#include "../asm_interrupt_handler.h" /* call_timer_int_handler(),
                                     call_keybd_int_handler() */

#include <timer_driver.h> /* init_timer() */
#include <keybd_driver.h> /* init_keybd() */
#include <install_handler.h>

/* Number of bits in a byte */
#define BYTE_LEN 8

/* Number of bytes that a trap gate occupies */
#define BYTES_PER_GATE 8

/* Mask for function handler address for upper bits */
#define OFFSET_UPPER_MASK 0xFFFF0000

/* Mask for function handler address for lower bits */
#define OFFSET_LOWER_MASK 0x0000FFFF

/* Trap gate flag masks */
#define PRESENT 0x00008000
#define DPL_0 0x00000000
#define D16 0x00000700
#define D32 0x00000F00
#define RESERVED_UPPER_MASK 0x0000000F

/* Handler installation error codes */
#define E_NO_INSTALL_KEYBOARD_HANDLER -2
#define E_NO_INSTALL_TIMER_HANDLER -3

/* Type of assembly wrapper for C interrupt handler functions */
typedef void asm_wrapper_t(void);

/*****
 *
 * Internal helper functions
 *
 *****/

/** @brief Installs an interrupt handler at idt_entry for timer and keyboard
 * interrupt handlers.
 *
 * If a tickback function pointer is provided, init_timer() will be called.
 * Else it is assumed that init_keybd() should be called instead, since there
 * are only 2 possible idt_entry values to accept.
 *
 * Furthermore, when casting from pointer types to integer types to pack
 * bits into the trap gate data structure, the pointer is first cast to
 * unsigned long then unsigned int, therefore there is no room for undefined
 * behavior.
 *
 * @param idt_entry Index into IDT table.
 * @param asm_wrapper Assembly wrapper to call C interrupt handler
 */

/* @param tickback Application provided callback function for timer interrupts.
 * @return 0 on success, -1 on error.
 */
int handler_install_in_idt(int idt_entry, asm_wrapper_t *asm_wrapper,
                          void (*tickback)(unsigned int)) {

    if (asm_wrapper == NULL) return -1;

    /* Only when installing the timer handler do we have a non-NULL tickback */
    if (tickback != NULL) {
        init_timer(tickback);
    } else {
        init_keybd();
    }

    /* Get address of trap gate for timer */
    void *idt_base_addr = idt_base();
    void *idt_entry_addr = idt_base_addr + (idt_entry * BYTES_PER_GATE);

    /* Exact offsets of each 32-bit word in the trap gate */
    unsigned int *idt_entry_addr_lower = idt_entry_addr;
    unsigned int *idt_entry_addr_upper = idt_entry_addr + (BYTES_PER_GATE / 2);
    if (idt_entry_addr_lower == NULL || idt_entry_addr_upper == NULL) return -1;

    /* Construct data for upper 32-bit word with necessary flags */
    unsigned int data_upper = 0;
    unsigned int offset_upper =
        ((unsigned int) (unsigned long) asm_wrapper) & OFFSET_UPPER_MASK;
    data_upper = offset_upper | PRESENT | DPL_0 | D32;

    /* Zero all bits that are not reserved, pack data into upper 32-bit word */
    *idt_entry_addr_upper = *idt_entry_addr_upper & RESERVED_UPPER_MASK;
    *idt_entry_addr_upper = *idt_entry_addr_upper | data_upper;

    /* Construct data for lower 32-bit word with necessary flags */
    unsigned int data_lower = 0;
    unsigned int offset_lower =
        ((unsigned int) (unsigned long) asm_wrapper) & OFFSET_LOWER_MASK;
    data_lower = (SEGSEL_KERNEL_CS << (2 * BYTE_LEN)) | offset_lower;

    /* Pack data into lower 32-bit word */
    *idt_entry_addr_lower = data_lower;

    /* This should always be the case after writing to IDT */
    assert(*idt_entry_addr_lower == data_lower);
    assert(*idt_entry_addr_upper == data_upper);

    return 0;
}

/*****
 *
 * Interface for device-driver initialization and timer callback
 *
 *****/

/** @brief The driver-library initialization function
 *
 * Installs the timer and keyboard interrupt handler.
 * NOTE: handler_install should ONLY install and activate the
 * handlers; any application-specific initialization should
 * take place elsewhere.
 * --
 * After installing both the timer and keyboard handler successfully,
 * interrupts are enabled.
 * --
 * @param tickback Pointer to clock-tick callback function
 * @return A negative error code on error, or 0 on success
 */

```

```
*/  
int handler_install(void (*tickback)(unsigned int)) {  
  
    /* While interrupt handlers are not set up, disable interrupts */  
    disable_interrupts();  
  
    /* Initialize and install timer handler */  
    int res = handler_install_in_idt(TIMER_IDT_ENTRY, call_timer_int_handler,  
                                     tickback);  
    if (res < 0) return E_NO_INSTALL_TIMER_HANDLER;  
  
    /* Initialize and install keyboard handler */  
    res = handler_install_in_idt(KEY_IDT_ENTRY, call_keybd_int_handler, NULL);  
    if (res < 0) return E_NO_INSTALL_KEYBOARD_HANDLER;  
  
    /* Interrupt handlers successfully installed, enable interrupts */  
    enable_interrupts();  
    return 0;  
}
```

```
.globl iret_travel

# iret_travel (ss, esp, eflags, cs, eip )
iret_travel:
    subl $4, %esp    # Point esp to last argument
    iret             # Consume all arguments and go to user mode!
```

./kern/kernel.c

```

/** @file kernel.c
 * @brief An initial kernel.c
 *
 * You should initialize things in kernel_main(),
 * and then run stuff.
 *
 * @author Harry Q. Bovik (hqbovik)
 * @author Fred Hacker (fhacker)
 * @bug No known bugs.
 */

#include <install_handler.h> /* handler_install() */
#include <common_kern.h>

/* libc includes. */
#include <stdio.h>
#include <simics.h> /* lprintf() */

/* multiboot header file */
#include <multiboot.h> /* boot_info */

/* x86 specific includes */
#include <x86/asm.h> /* enable_interrupts() */

#include <x86/cr.h> /* get_cr3() */

#include <exec2obj.h> /* MAX_EXECNAME_LEN */

#include <loader.h> /* execute_user_program() */
#include <console.h> /* clear_console(), putbytes() */
#include <keybd_driver.h> /* readline() */

volatile static int __kernel_all_done = 0;

/* Think about where this declaration
 * should be... probably not here!
 */
void tick(unsigned int numTicks) {
    //lprintf("numTicks: %d\n", numTicks);
}

/** @brief Kernel entrypoint.
 *
 * This is the entrypoint for the kernel.
 *
 * @return Does not return
 */
int
kernel_main( mbinfo_t *mbinfo, int argc, char **argv, char **envp )
{
    /* FIXME: What to do with mbinfo and envp? */
    (void)mbinfo;
    (void)envp;

    /* initialize device-driver library */
    handler_install(tick);

    clear_console();

    /* TODO
     * When kernel_main() begins, interrupts are DISABLED.
     * You should delete this comment, and enable them --
     * when you are ready.
     */

    lprintf( "Hello from a brand new kernel!" );
    putbytes("executable user programs:\n", 26);

```

```

    putbytes("loader_test1\n", 13);
    putbytes("loader_test2\n", 13);
    putbytes("getpid_test1\n", 13);

    /** On kernel_main() entry, all control registers are 0 */
    //lprintf("cr1: %p", (void *) get_cr3());
    //lprintf("cr2: %p", (void *) get_cr3());
    //lprintf("cr3: %p", (void *) get_cr3());
    //lprintf("cr4: %p", (void *) get_cr3());
    //char * nullp = 0;
    //lprintf("garbage at address 0x0:%d", *nullp);
    //lprintf("&nullp:%p", &nullp);

    while (!__kernel_all_done) {
        int n = MAX_EXECNAME_LEN;
        char s[n];

        /* Display prompt */
        putbytes("pebbles>", 8);
        int res = readline(s, n);

        if (res == n)
            continue; /* Executable name too large */

        /* Swap \n returned by readline for null-terminator */
        s[res - 1] = '\0';

        lprintf("Executing: %s", s);

        char *user_argv = (char *)s;
        execute_user_program(s, 1, &user_argv);
    }

    return 0;
}

```

./kern/keybd_driver.c

```

/** @file keybd_driver.c
 * @brief Contains functions that help the user type into the console
 *
 * @bug No known bugs.
 *
 * Since unbounded arrays are used, let 'size' be the number of elements in the
 * array that we care about, and 'limit' be the actual length of the array.
 * Whenever the size of the array == its limit, the array limit is doubled.
 * Doubling is only allowed if the current limit <= UINT32_MAX to prevent
 * overflow.
 *
 * Indexing into circular arrays is just done modulo the limit of the array.
 *
 * The two functions in the keyboard driver interface readchar() amd readline()
 * are closely connected to one another. The specification for readline()
 * states that "Characters not placed into the specified buffer should remain
 * available for other calls to readline() and/or readchar()." To put it
 * concisely, we have the implication:
 *
 * char not committed to any buffer => char available for other calls
 *
 * i.e.
 *
 * char not available for other calls => char committed to some buffer
 *
 * The question now is under what circumstance do we promise that a char
 * is not available? It is reasonable to conclude that the above implication
 * is in fact a bi-implication. Therefore:
 *
 * char not available for other calls <=> char committed to some buffer
 *
 * Now since readchar() always reads the next character in the keyboard buffer,
 * it is then conceivable that if a readchar() not called by readline() takes
 * the next character off the keyboard buffer, then readline() would "skip"
 * a character. But then, the specification says that:
 *
 * "Since we are operating in a single-threaded environment, only one of
 * readline() or readchar() can be executing at any given point."
 *
 * Therefore we will never have readline() and readchar() concurrently
 * executing in separate threads in the context of the same process, since
 * when we speak of threads, we refer to threads in the same process.
 *
 * Therefore we need not worry about the case where a readchar() that is
 * not invoked by readline() is called in the middle of another call to
 * readline().
 *
 * @author Nicklaus Choo (nchoo)
 * @bug No known bugs
 */

#include <x86/keyhelp.h> /* process_scancode() */
#include <x86/video_defines.h> /* CONSOLE_HEIGHT, CONSOLE_WIDTH */
#include <malloc.h> /* calloc */
#include <stddef.h> /* NULL */
#include <assert.h> /* assert() */
#include <console.h> /* putbyte() */
#include <string.h> /* memcpy() */
#include <x86/asm.h> /* process_scancode() */
#include <x86/interrupt_defines.h> /* INT_CTL_PORT */
#include <ctype.h> /* isprint() */
#include <./console_driver.h> /* _putbyte() */
#include <./keybd_driver.h> /* uba */

/* Keyboard buffer */
static uba *key_buf = NULL;

```

```

int readchar(void);

/***** Internal helper functions *****/
/*
 * @brief Checks invariants for unbounded arrays
 *
 * The invariants here are implemented as asserts so in the even that
 * an assertion fails, the developer knows exactly which assertion
 * fails. There are many invariants since on top of being an
 * unbounded array, it is also a circular array.
 *
 * @param arr Pointer to uba to be checked
 * @return 1 if valid uba pointer, 0 otherwise
 */
int is_uba(uba *arr) {

    /* non-NULL check */
    assert(arr != NULL);

    /* size check */
    assert(0 <= arr->size && arr->size < arr->limit);

    /* limit check */
    assert(0 < arr->limit);

    /* first and last check */
    if (arr->first <= arr->last) {
        assert(arr->size == arr->last - arr->first);
    } else {
        assert(arr->size == arr->limit - arr->first + arr->last);
    }

    /* first and last in bounds check */
    assert(0 <= arr->first && arr->first < arr->limit);
    assert(0 <= arr->last && arr->last < arr->limit);

    /* data non-NULL check */
    assert(arr->data != NULL);
    return 1;
}

/** @brief Initializes a uba and returns a pointer to it
 *
 * Fatal errors are only thrown if size == 0 or size == 1 and yet
 * init_uba() returns NULL (out of heap space)
 *
 * @param type Element type in the uba
 * @param limit Actual size of the uba.
 * @return Pointer to valid uba if successful, NULL if malloc() or
 *         calloc() fails or if limit < 0
 */
uba *uba_new(int limit) {
    if (limit <= 0) return NULL;

    /* Convert limit to unsigned */
    uint32_t _limit = (uint32_t) limit;
    assert(0 <= _limit);

    /* Allocate memory for uba struct */
    uba *new_ubap = malloc(sizeof(uba));
    assert(new_ubap != NULL);
    if (new_ubap == NULL) return NULL;

```

./kern/keyboard_driver.c

```

/* Allocate memory for the array */
void *data = NULL;
data = calloc(_limit, sizeof(uint8_t));
assert(data != NULL);
if (data == NULL) return NULL;

/* Set fields of unbounded array uba */
new_ubap->size = 0;
new_ubap->limit = _limit;
new_ubap->first = 0;
new_ubap->last = 0;
new_ubap->data = data;

return new_ubap;
}

/** @brief Resizes the unbounded array
 *
 * @param arr Pointer to a uba
 * @param Pointer to the original uba if no resize needed, bigger uba o/w.
 */
uba *uba_resize(uba *arr) {
    assert(is_uba(arr));
    if (arr->size == arr->limit) {

        /* Only resize if sufficient */
        if (arr->size < UINT32_MAX / 2) {
            uba *new_arr = uba_new(arr->size * 2);
            assert(is_uba(new_arr));

            /* Copy elements over */
            while(arr->size > 0) {
                uint8_t elem = uba_rem(arr);
                uba_add(new_arr, elem);
            }
            /* Free the old array */
            free(arr->data);
            free(arr);

            /* Return new array */
            return new_arr;
        }
        /* at limit but cannot resize, OK for now but error will be thrown
        * if want to add a new element to arr
        */
    }
    return arr;
}

/** @brief Adds new element to the array and resizes if needed
 *
 * @param arr Unbounded array pointer we want to add to
 * @param elem Element to add to the unbounded array
 * @return Void.
 */
void uba_add(uba *arr, uint8_t elem) {
    assert(is_uba(arr));
    assert(arr->size < arr->limit); /* Total memory is 256 MiB and since
                                   each array element is 1 byte we
                                   will never reach max possible limit
                                   of 2^32 */

    /* Insert at index one after last element in array */
    arr->data[arr->last] = elem;

    /* Update last index and increment size*/
    arr->last = (arr->last + 1) % arr->limit;

    arr->size += 1;
    assert(is_uba(arr));

    /* Resize if necessary */
    arr = uba_resize(arr);
    assert(is_uba(arr));
}

/** @brief Removes first character of the uba
 *
 * @param arr Pointer to uba
 * @return First character of the uba
 */
uint8_t uba_rem(uba *arr) {
    assert(is_uba(arr));

    /* Get first element and 'remove' from the array, decrement size */
    uint8_t elem = arr->data[arr->first];
    arr->first = (arr->first + 1) % arr->limit;
    arr->size -= 1;
    assert(is_uba(arr));

    /* Return 'popped off' element */
    return elem;
}

/** @brief Checks if uba is empty
 *
 * @param arr Pointer to uba
 * @return 1 if empty, 0 otherwise.
 */
int uba_empty(uba *arr) {
    assert(is_uba(arr));
    return arr->size == 0;
}

/** @brief Interrupt handler which reads in raw bytes from keystrokes. Reads
 *
 * incoming bytes to the keyboard buffer key_buf, which has an
 * amortized constant time complexity for adding elements. So it
 * returns quickly
 *
 * @return Void.
 */
void keyboard_int_handler(void) {

    /* Read raw byte and put into raw character buffer */
    uint8_t raw_byte = inb(KEYBOARD_PORT);
    uba_add(key_buf, raw_byte);

    /* Acknowledge interrupt and return */
    outb(INT_CTL_PORT, INT_ACK_CURRENT);
}

/** @brief Initialize the keyboard interrupt handler and associated data
 *
 * structures
 *
 * Memory for keybd_buf is allocated here.
 *
 * @return Void.
 */
void init_keybd(void) {

    /* Initialize the raw_byte buffer */
    key_buf = uba_new(CONSOLE_HEIGHT * CONSOLE_WIDTH);
    assert(key_buf != NULL);
}

/** @brief Keeps calling readchar() until another valid char is read and
 *
 * returns it.

```

./kern/keyboard_driver.c

```

*
* @return A valid character when readchar() doesn't return -1.
*/
char get_next_char(void) {

    /* Get the next char value off the keyboard buffer */
    int res;
    while((res = readchar()) == -1) continue;
    assert(res >= 0);

    /* Tricky type conversions to avoid undefined behavior */
    char char_value = (uint8_t) (unsigned int) res;
    return char_value;
}

/*****
*/
/* Keyboard driver interface */
/*
*/
/*****

/** @brief Returns the next character in the keyboard buffer
*
* This function does not block if there are no characters in the keyboard
* buffer
* --
* No other process will call readchar() concurrently with readline() since
* we only have 1 kernal process running and have a single thread.
* --
* @return The next character in the keyboard buffer, or -1 if the keyboard
*         buffer is currently empty
**/
int readchar(void) {

    assert(key_buf != NULL);

    /* uba invariants are checked whenever we access the data structure, and so
    * interrupts are disabled during the entire call in order to prevent
    * changes to the data structure at the start of and at the end of a call.
    */
    disable_interrupts();
    if (uba_empty(key_buf)) {
        enable_interrupts();
        return -1;
    }
    raw_byte next_byte = uba_rem(key_buf);
    enable_interrupts();

    /* Get augmented character */
    aug_char next_char = process_scancode(next_byte);

    /* Get simplified character */
    if (KH_HASDATA(next_char)) {
        if (KH_ISMAKE(next_char)) {
            unsigned char next_char_value = KH_GETCHAR(next_char);
            return (int) (unsigned int) next_char_value;
        }
    }
    return -1;
}

/** @brief Reads a line of characters into a specified buffer
*
* If the keyboard buffer does not already contain a line of input,
* readline() will spin until a line of input becomes available.
*
* If the line is smaller than the buffer, then the complete line,

```

```

* including the newline character, is copied into the buffer.
*
* If the length of the line exceeds the length of the buffer, only
* len characters should be copied into buf.
*
* Available characters should not be committed into buf until
* there is a newline character available, so the user has a
* chance to backspace over typing mistakes.
*
* While a readline() call is active, the user should receive
* ongoing visual feedback in response to typing, so that it
* is clear to the user what text line will be returned by
* readline().
* --
* the definition of a line in readline() is different from a row. A
* carriage-return will return the cursor to its initial position at the
* start of the call to readline(). Backspaces will always work and only
* do nothing if the cursor is at the initial position at the start of the
* call to readline.
*
* Since it is only meaningful that the user can see exactly what was written
* to buf, If len is valid the moment the user types len bytes readline() will
* return. We prevent the user from typing more than len characters into the
* console.
* --
* @param buf Starting address of buffer to fill with a text line
* @param len Length of the buffer
* @return The number of characters in the line buffer,
*         or -1 if len is invalid or unreasonably large.
*/
int readline(char *buf, int len) {

    /* buf == NULL so invalid buf */
    if (buf == NULL) return -1;

    /* len < 0 so invalid len */
    if (len < 0) return -1;

    /* len == 0 so no need to copy */
    if (len == 0) return 0;

    /* get original cursor position for start of line relative to scroll */
    int start_row, start_col;
    get_cursor(&start_row, &start_col);

    /* Allocate space for temporary buffer */
    char temp_buf[len];

    /* Initialize index into temp_buf */
    int i = 0;
    int written = 0; /* characters written so far */
    char ch;

    while ((ch = get_next_char()) != '\n' && written < len) {

        /* ch, i, written is always in range */
        assert(0 <= i && i < len);
        assert(0 <= written && written < len);

        /* If at front of buffer, Delete the character if backspace */
        if (ch == '\b') {

            /* If at start_row, start_col, do nothing as don't delete prompt */
            int row, col;
            get_cursor(&row, &col);
            assert(row * CONSOLE_WIDTH + col >= start_row * CONSOLE_WIDTH + start_col);
            if (!(row == start_row && col == start_col)) {

```

```
    assert(i > 0);

    /* Print to screen and update initial cursor position if needed*/
    scrolled_putbyte(ch, &start_row, &start_col);

    /* update i and buffer */
    i--;
    temp_buf[i] = ' ';
}
/* '\r' sets cursor to position at start of call. Don't overwrite prompt */
} else if (ch == '\r') {

    /* Set cursor to start of line w.r.t start of call, i to buffer start */
    set_cursor(start_row, start_col);
    i = 0;

    /* Regular characters just write, unprintables do nothing */
} else {

    /* print on screen and update initial cursor position if needed */
    scrolled_putbyte(ch, &start_row, &start_col);

    /* write to buffer */
    if (isprint(ch)) {
        temp_buf[i] = ch;
        i++;
        if (i > written) written = i;
    }
}
}
assert(written <= len);
if (ch == '\n') {

    /* Only write the newline if there's space for it in the buffer */
    if (written < len) {
        putbyte(ch);
        temp_buf[i] = '\n';
        i++;
        if (i > written) written = i;
    }
} else {
    assert(written == len);
}
memcpy(buf, temp_buf, written);
return written;
}
```


./kern/loader.c

```

/**
 * The 15-410 kernel project.
 * @name loader.c
 *
 * Functions for the loading
 * of user programs from binary
 * files should be written in
 * this file. The function
 * elf_load_helper() is provided
 * for your use.
 *
 * The loader should never interact directly with
 * virtual memory. Rather it should call functions
 * defined in the process manager module (which itself
 * will be responsible for talking to the VM module).
 */
/*{*/

/* --- Includes --- */
#include <loader.h>
#include <page.h>      /* PAGE_SIZE */
#include <string.h>     /* strcmp, memcpy */
#include <exec2obj.h>   /* exec2obj_TOC */
#include <elf_410.h>    /* simple_elf_t, elf_load_helper */
#include <limits.h>     /* UINT_MAX */
#include <task_manager.h> /* task_new, task_prepare, task_set */
#include <memory_manager.h> /* {disable,enable}_write_protection */

#include <simics.h> /* lprintf */

/* --- Local function prototypes --- */

/* TODO: Move this to a helper file.
 *
 * Having a helper means we evaluate the arguments before expanding _MIN
 * and therefore avoid evaluating A and B multiple times. */
#define _MIN(A, B) ((A) < (B) ? (A) : (B))
#define MIN(A,B) _MIN(A,B)

/** Copies data from a file into a buffer.
 *
 * @param filename the name of the file to copy data from
 * @param offset the location in the file to begin copying from
 * @param size the number of bytes to be copied
 * @param buf the buffer to copy the data into
 *
 * @return number of bytes copied on success. Negative value on failure.
 */
int
getbytes( const char *filename, int offset, int size, char *buf )
{
    if (!filename || !buf || offset < 0 || size < 0) {
        lprintf("Loader [getbytes]: Invalid arguments");
        return -1;
    }

    /* Find file in TOC */
    int i;
    for (i=0; i < exec2obj_userapp_count; ++i) {
        if (strcmp(filename, exec2obj_userapp_TOC[i].execname, MAX_EXECNAME_LEN) ==
            0) {
            break;
        }
    }

    if (i == exec2obj_userapp_count) {
        lprintf("Loader [getbytes]: Executable not found");
        return -1;
    }

    /* FIXME: Spec is unclear. Should we copy as much as we can, or should
     * only copy if there are enough bytes in the executable? */
    /*if (offset + size >= exec2obj_userapp_TOC[i].execlen)
     * return -1; /* Asking for more bytes than are available */

    int bytes_to_copy = MIN(size, exec2obj_userapp_TOC[i].execlen - offset);

    memcpy(buf, exec2obj_userapp_TOC[i].execbytes + offset, bytes_to_copy);

    return bytes_to_copy;
}

/** @brief Transplants program data into virtual memory.
 * Assumes paging is enabled and that virtual memory
 * has been setup.
 *
 * @arg se_hdr Elf header
 *
 * @return 0 on success, negative value on failure.
 */
static int
transplant_program_memory( simple_elf_t *se_hdr )
{
    /* Disable write-protection temporarily so we may
     * copy data into read-only regions. */
    disable_write_protection();

    /* FIXME: This error checking is kinda hacky
     * int i = 0;

     * TODO: Zero out bytes between memory regions */

    /* We rely on the fact that virtual-memory is
     * enabled to "transplant" program data. Notice
     * that this is only possible because program data is
     * resident on kernel memory which is direct-mapped. */
    i += getbytes(se_hdr->e_fname,
        (unsigned int) se_hdr->e_textoff,
        (unsigned int) se_hdr->e_textlen,
        (char *) se_hdr->e_textstart);

    i += getbytes(se_hdr->e_fname,
        (unsigned int) se_hdr->e_dataoff,
        (unsigned int) se_hdr->e_datlen,
        (char *) se_hdr->e_datastart);

    i += getbytes(se_hdr->e_fname,
        (unsigned int) se_hdr->e_rodataoff,
        (unsigned int) se_hdr->e_rodatlen,
        (char *) se_hdr->e_rodatastart);

    /* Re-enable write-protection bit. */
    enable_write_protection();

    return i;
}

/** @brief Puts arguments on stack with format required by _main entrypoint.
 *
 * This entrypoint is defined in 410user/crt0.c and is used by all user
 * programs.
 */

```

```
static uint32_t *
configure_stack( int argc, char **argv )
{
    /* TODO: In the future, when "receiver" function is implemented, loader
     * should also add entry point, user registers and data segment selectors
     * on the stack. For registers, just initialize most to 0 or something. */

    uint32_t *esp = (uint32_t *)UINT_MAX;

    *esp = argc;

    if (argc == 0) {
        *--esp = 0;
        return esp;
    }

    esp -= argc; /* sizeof(char *) == sizeof(uint32_t *) */
    memcpy(esp, argv, argc * sizeof(char *)); /* Put argv on stack */
    esp--;

    *esp-- = UINT_MAX; /* Put stack_high on stack */
    *esp = UINT_MAX - PAGE_SIZE; /* Put stack_low on stack */

    /* Functions expect esp to point to return address on entry.
     * Therefore we just point it to some garbage, since _main
     * is never supposed to return. */
    esp--;

    return esp;
}

/** @brief Run a user program indicated by filename.
 *     Assumes virtual memory module has been initialized.
 *
 * @arg fname Name of program to run.
 *
 * @return 0 on success, negative value on error.
 */
int
execute_user_program( const char *fname, int argc, char **argv )
{
    /* Load user program information */
    simple_elf_t se_hdr;
    if (elf_check_header(fname) == ELF_NOTELF)
        return -1;

    if (elf_load_helper(&se_hdr, fname) == ELF_NOTELF)
        return -1;

    lprintf("creating task");

    /* FIXME: Hard coded pid and tid for now */
    if (task_new(0, 0, &se_hdr) < 0)
        return -1;

    lprintf("preparing task");

    /* Enable VM */
    if (task_prepare(0) < 0)
        return -1;

    lprintf("transplanting memory task");

    if (transplant_program_memory(&se_hdr) < 0)
        return -1;
}
```

```
    lprintf("configuring stack");

    uint32_t *esp = configure_stack(argc, argv);

    lprintf("setting task");

    task_set(0, (uint32_t)esp, se_hdr.e_entry);

    return 0;
}

/* we try to use the physical addresses */
/*@{*/
```

```
#include <stddef.h>
#include <malloc.h>
#include <malloc_internal.h> /* _malloc family of functions */

/* safe versions of malloc functions */
void *malloc(size_t size)
{
    return _malloc(size);
}

void *memalign(size_t alignment, size_t size)
{
    return _memalign(alignment, size);
}

void *calloc(size_t nelt, size_t eltsize)
{
    return _calloc(nelt, eltsize);
}

void *realloc(void *buf, size_t new_size)
{
    return _realloc(buf, new_size);
}

void free(void *buf)
{
    _free(buf);
}

void *smallloc(size_t size)
{
    return _smallloc(size);
}

void *smemalign(size_t alignment, size_t size)
{
    return smemalign(alignment, size);
}

void sfree(void *buf, size_t size)
{
    sfree(buf, size);
}
```

./kern/memory_manager.c

```

/** Virtual memory manager
 *
 * */

#include <memory_manager.h>
#include <stdint.h> /* uint32_t */
#include <stddef.h> /* NULL */
#include <malloc.h> /* smemalign, sfree */
#include <elf_410.h> /* simple_elf_t */
#include <assert.h> /* assert, affirm */
#include <page.h> /* PAGE_SIZE */
#include <x86/cr.h> /* {get,set}_cr0 */
#include <string.h> /* memset, memcpy */
#include <common_kern.h> /* USER_MEM_START */

#include <simics.h> /* lprintf */

#define PAGING_FLAG (1 << 31)
#define WRITE_PROTECT_FLAG (1 << 16)

#define PAGE_DIRECTORY_INDEX 0xFFC00000
#define PAGE_TABLE_INDEX 0x003FF000
#define PAGE_OFFSET 0x0000FFF

#define PAGE_DIRECTORY_SHIFT 22
#define PAGE_TABLE_SHIFT 12

/* Get page directory or page table index from a logical/linear address. */
#define PD_INDEX(addr) \
    ((PAGE_DIRECTORY_INDEX & (addr)) >> PAGE_DIRECTORY_SHIFT)
#define PT_INDEX(addr) \
    ((PAGE_TABLE_INDEX & (addr)) >> PAGE_TABLE_SHIFT)

/* Flags for page directory and page table entries */
#define PRESENT_FLAG 1 << 0
#define RW_FLAG 1 << 1
#define US_FLAG 1 << 2
#define GLOBAL_FLAG 1 << 8

#define PE_USER_READABLE (PRESENT_FLAG | US_FLAG)
#define PE_USER_WRITABLE (PE_USER_READABLE | RW_FLAG)

/* Set global flag so TLB doesn't flush kernel entries */
#define PE_KERN_READABLE (PRESENT_FLAG | GLOBAL_FLAG | RW_FLAG)
#define PE_KERN_WRITABLE (PE_KERN_READABLE | RW_FLAG)
#define PE_UNMAPPED 0

/* FIXME: Temporary variable for enabling allocation of physical frames.
 * Only to be used for user memory. Starts at USER_MEM_START, where
 * the first phys frames are available. */
static uint32_t next_free_phys_frame;

/** Whether page is read only or also writable. */
typedef enum write_mode write_mode_t;
enum write_mode { READ_ONLY, READ_WRITE };

static int get_next_free_frame( uint32_t *frame );
static uint32_t num_free_frames( void );
static uint32_t *get_pte( uint32_t **ptd, uint32_t virtual_address );
static void allocate_frame( uint32_t **ptd,
    uint32_t virtual_address, write_mode_t write_mode );
static int allocate_region( void *ptd, void *start,
    uint32_t len, write_mode_t write_mode );
static void enable_paging( void );
static void disable_paging( void );

/** Initialize virtual memory. */

```

```

int
vm_init( void )
{
    next_free_phys_frame = USER_MEM_START;

    assert((next_free_phys_frame & (PAGE_SIZE - 1)) == 0);

    return 0;
}

/** Allocate memory for new task at given page table directory.
 * Assumes page table directory is empty. Sets appropriate
 * read/write permissions. To copy memory over, set the WP flag
 * in the CR0 register to 0 - this will make it so that write
 * protection is ignored by the paging mechanism.
 *
 * Allocated pages are initialized to 0.
 *
 * TODO: Implement ZFOD here. (Handler should probably be defined
 * elsewhere, though)
 */
int
vm_task_new ( void *ptd, simple_elf_t *elf,
    uint32_t stack_lo, uint32_t stack_len )
{
    affirm(ptd);

    lprintf("Direct mapping kernel");

    /* Direct map all 16MB for kernel, setting correct permission bits */
    for (uint32_t addr = 0; addr < USER_MEM_START; addr += PAGE_SIZE) {
        uint32_t *pte = get_pte(ptd, addr);
        if (addr == 0) {
            *pte = addr | PE_UNMAPPED; /* Leave NULL unmapped. */
        } else {
            *pte = addr | PE_KERN_WRITABLE;
        }
    }

    lprintf("Direct mapping kernel");

    /* Allocate regions with appropriate read/write permissions.
     * TODO: Free allocated regions if later allocation fails. */
    int i = 0;

    lprintf("Allocating regions");
    i += allocate_region(ptd, (void *)elf->e_txtstart, elf->e_txtlen, READ_ONLY);
    lprintf("Allocating regions");
    i += allocate_region(ptd, (void *)elf->e_datstart, elf->e_datlen, READ_WRITE);
    lprintf("Allocating regions");
    i += allocate_region(ptd, (void *)elf->e_rodstart, elf->e_rodlen, READ_ONLY);
    lprintf("Allocating regions");
    i += allocate_region(ptd, (void *)elf->e_bssstart, elf->e_bsslen, READ_WRITE);
    lprintf("Allocating regions");
    i += allocate_region(ptd, (void *)stack_lo, stack_len, READ_WRITE);
    lprintf("Allocated all regions");

    return i;
}

/** @brief Sets new page table directory and enables paging. */
void
vm_enable_task( void *ptd )
{
    uint32_t cr3 = get_cr3();
    /* Unset top 20 bits where new page table will be stored.*/

```

./kern/memory_manager.c

```

    cr3 &= PAGE_SIZE - 1;
    cr3 |= (uint32_t)ptd;

    set_cr3(cr3);
}

/** @brief Enables write_protect flag in cr0, allowing
 * kernel to bypass VM's read-only protection. */
void
enable_write_protection( void )
{
    uint32_t current_cr0 = get_cr0();
    set_cr0(current_cr0 | WRITE_PROTECT_FLAG);
}

/** @brief Disables write_protect flag in cr0, stopping
 * kernel from bypassing VM's read-only protection. */
void
disable_write_protection( void )
{
    uint32_t current_cr0 = get_cr0();
    set_cr0(current_cr0 & (~WRITE_PROTECT_FLAG));
}

/** Allocate new pages in a given process' virtual memory. */
int
vm_new_pages ( void *ptd, void *base, int len )
{
    // TODO: Implement
    (void)ptd;
    (void)base;
    (void)len;
    return -1;
}

/* ----- HELPER FUNCTIONS ----- */

/** Gets new page-aligned physical frame.
 *
 * @arg frame Memory location in which to store new frame address
 *
 * @return 0 on success, negative value on failure
 * */
static int
get_next_free_frame( uint32_t *frame )
{
    uint32_t free_frame = next_free_phys_frame;

    if (num_free_frames() == 0)
        return -1;

    next_free_phys_frame += PAGE_SIZE;

    assert((free_frame & (PAGE_SIZE - 1)) == 0);

    *frame = free_frame;
    return 0;
}

/** Gets number of remaining free physical frames. */
static uint32_t
num_free_frames( void )
{
    int remaining = machine_phys_frames() - (next_free_phys_frame / PAGE_SIZE);
    affirm(remaining >= 0); /* We should never allocate memory we don't have! */
    return (uint32_t)remaining;
}

/** Gets pointer to page table entry in a given page directory.
 *
 * Allocates page table if necessary. */
static uint32_t *
get_pte( uint32_t **ptd, uint32_t virtual_address )
{
    affirm(ptd);
    uint32_t pd_index = PD_INDEX(virtual_address);
    uint32_t pt_index = PT_INDEX(virtual_address);

    if (!(uint32_t)ptd[pd_index] & PRESENT_FLAG) {
        /* Allocate new page table, which must be page-aligned */
        ptd[pd_index] = smemalign(PAGE_SIZE, PAGE_SIZE);
        affirm(ptd[pd_index]);

        /* Initialize all page table entries as non-present */
        memset(ptd[pd_index], 0, PAGE_SIZE);

        /* Set all page directory entries as writable, determine
         * whether truly writable in page table entry. */
        ptd[pd_index] = (uint32_t *)((uint32_t)ptd[pd_index] | PE_USER_WRITABLE);
    }

    uint32_t *page_table = ptd[pd_index];

    return page_table + pt_index;
}

/** Allocate new frame at given virtual memory address.
 *
 * Allocates page tables on demand.
 *
 * If memory location already had a frame, this crashes.
 * */
static void
allocate_frame( uint32_t **ptd, uint32_t virtual_address, write_mode_t write_mode )
{
    affirm(ptd);

    uint32_t free_frame;
    affirm(get_next_free_frame(&free_frame) == 0);

    uint32_t *pte = get_pte(ptd, virtual_address);
    uint32_t pt_index = PT_INDEX(virtual_address);
    affirm(((uint32_t)ptd[pt_index] & PRESENT_FLAG) == 0); /* Ensure unallocated */

    *pte = free_frame;

    /* FIXME: Hack for until we implement ZFOD. Do we even want to guarantee
     * zero-filled pages for the initially allocated regions? Seems like
     * .bss and new_pages are the only ones required to be zeroed out by spec.*/
    /* ATOMICALLY start */
    disable_paging();
    memset((void *)free_frame, 0, PAGE_SIZE);
    enable_paging();
    /* ATOMICALLY end*/

    if (write_mode == READ_WRITE)
        *pte |= PE_USER_WRITABLE;
    else
        *pte |= PE_USER_READABLE;
}

/** Allocates a memory region in virtual memory.
 *
 * If there aren't enough physical frames to satisfy allocation
 * request, region is not allocated and function returns a negative
 * value.

```

```
*
* @arg ptd      Pointer to page directory
* @arg start    Virtual memory address for start of region to be allocated
* @arg len      Length of region to be allocated
* @arg write_mode 0 if read-only region, non-zero value if writable
*
* @return 0 on success, negative value on failure.
* */
static int
allocate_region( void *ptd, void *start, uint32_t len, write_mode_t write_mode )
{
    /* Ensure we have enough free frames to fulfill request */
    if (num_free_frames() < (len + PAGE_SIZE - 1) / PAGE_SIZE)
        return -1;

    uint32_t curr = (uint32_t)start;

    /* Allocate 1 frame at a time. */
    while (curr < (uint32_t)start + len) {
        allocate_frame((uint32_t **)ptd, curr, write_mode);
        curr += PAGE_SIZE;
    }

    return 0;
}

/** @brief Enables paging mechanism. */
static void
enable_paging( void )
{
    uint32_t current_cr0 = get_cr0();
    set_cr0(current_cr0 | PAGING_FLAG);
}

/** @brief Disables paging mechanism. */
static void
disable_paging( void )
{
    uint32_t current_cr0 = get_cr0();
    set_cr0(current_cr0 & (~PAGING_FLAG));
}
```

./kern/task_manager.c

```

/** @brief Module for management of tasks.
 * Includes context switch facilities. */

#include <task_manager.h>
#include <limits.h> /* UINT_MAX */
#include <eflags.h> /* get_eflags*/
#include <seg.h> /* SEGSEL_... */
#include <stdint.h> /* uint32_t */
#include <stddef.h> /* NULL */
#include <malloc.h> /* malloc, smemalign, free, sfree */
#include <elf_410.h> /* simple_elf_t */
#include <page.h> /* PAGE_SIZE */
#include <string.h> /* memset */
#include <assert.h> /* affirm, assert */
#include <simics.h> /* sim_reg_process */
#include <iret_travel.h> /* iret_travel */
#include <memory_manager.h> /* vm_task_new, vm_enable_task */

/** @brief Pointer to first task control block. */
pcb_t *pcb_list_start = NULL;

static int find_pcb( int pid, pcb_t **pcb );
static int new_pcb( int pid );

static int find_tcb( int tid, tcb_t **pcb );
static int new_tcb( int pid, int tid );

static uint32_t get_user_eflags( void );

/* Create new task
 *
 * @arg pid Task id for new task
 * @arg tid Thread id for new thread
 * @arg elf Elf header for use in allocating new task's memory
 *
 * @return 0 on success, negative value on failure.
 * */
int
task_new( int pid, int tid, simple_elf_t *elf )
{
    // TODO: Think about preconditions for this.
    // Paging fine, how about making it a critical section?

    printf("creating pcb");

    if (new_pcb(pid) < 0)
        return -1;

    printf("creating tcb");

    /* TODO: Deallocate pcb if this fails */
    if (new_tcb(pid, tid) < 0)
        return -1;

    /* Allocate VM. Stack currently starts at top most address
     * and is PAGE_SIZE long. */
    pcb_t *pcb;
    printf("finding pcb ");
    affirm(find_pcb(pid, &pcb) == 0);
    printf("creating vm task_new ");
    vm_task_new(pcb->ptd, elf, UINT_MAX, PAGE_SIZE);

    printf("registering process w/ simics");
#ifdef NDEBUG
    /* Register this task with simics for better debugging */
    sim_reg_process(pcb->ptd, elf->e_fname);
#endif
}

#endif

return 0;
}

/** NOTE: Not to be used in context-switch, only when running task
 * for the first time
 *
 * Enables virtual memory of task. Use this before transplanting data
 * into task's memory.
 * */
int
task_prepare( int pid )
{
    /* Likely messing up direct mapping of kernel memory, and
     * some instruction after task_prepare is being seen as invalid?*/
    pcb_t *pcb;
    if (find_pcb(pid, &pcb) < 0)
        return -1;

    printf("before enabling vm");
    /* Enable VM */
    vm_enable_task(pcb->ptd);
    printf("after enabling vm");

    return 0;
}

/** NOTE: Not to be used in context-switch, only when running task
 * for the first time
 *
 * Should only ever be called once, and after task has been initialized
 * after a call to new_task.
 * The caller is supposed to install memory on the new task before
 * calling this function. Stack pointer should be appropriately set
 * if any arguments have been loaded on stack.
 *
 * @arg tid Id of thread to run
 * @arg esp Stack pointer
 * @arg entry_point First program instruction
 *
 * @return Never returns.
 * */
void
task_set( int tid, uint32_t esp, uint32_t entry_point )
{
    tcb_t *tcb;
    affirm(find_tcb(tid, &tcb) == 0);
    pcb_t *pcb = tcb->owning_task;

    if (!pcb->prepared) {
        task_prepare(pcb->pid);
    }

    printf("before iret travel");

    /* We're currently going directly to entry point. In the future,
     * however, we should go to some "receiver" function which appropriately
     * sets user registers and segment selectors, and lastly RETs to
     * the entry_point. */
    iret_travel(SEGSEL_USER_DS, esp, get_user_eflags(),
                SEGSEL_USER_CS, entry_point);

    /* NOTREACHED */
    affirm(0);
}

```

./kern/task_manager.c

```

/* Aka context_switch */
void
task_switch( int pid )
{
    (void)pid;
    /* TODO: Unimplemented */
}

/* ----- HELPER FUNCTIONS ----- */

/** @brief Returns eflags with PL altered to 3 */
static uint32_t
get_user_eflags( void )
{
    /* Any IOPL | EFL_IOPL_RING3 == EFL_IOPL_RING3 */
    return get_eflags() | EFL_IOPL_RING3;
}

/** Looks for pcb with given pid.
 *
 * @arg pid Task id to look for
 * @arg pcb Memory location where to store pcb, if found.
 *
 * @return 0 on success, negative value on error. */
static int
find_pcb( int pid, pcb_t **pcb )
{
    // TODO: Actually implement the search
    if (!pcb_list_start)
        return -1;
    *pcb = pcb_list_start;
    return 0;
}

/** Looks for tcb with given tid.
 *
 * @arg tid Thread id to look for
 * @arg tcb Memory location where to store tcb, if found.
 *
 * @return 0 on success, negative value on error. */
static int
find_tcb( int tid, tcb_t **tcb )
{
    // TODO: Actually implement the search
    if (!pcb_list_start)
        return -1;
    *tcb = pcb_list_start->first_thread;
    return 0;
}

/* Initializes new pcb.
 *
 * TODO: Should we initialize a TCB here as well?
 * Does it make sense for a task with no threads to exist? */
static int
new_pcb( int pid )
{
    lprintf("smemalign");
    /* Ensure alignment of page table directory */
    void *ptd = smemalign(PAGE_SIZE, PAGE_SIZE);
    lprintf("after smemalign");
    if (!ptd)
        return -1;

    lprintf("assert");

```

```

    assert(((uint32_t)ptd & (PAGE_SIZE - 1)) == 0);

    lprintf("malloc");
    pcb_t *pcb = malloc(sizeof(pcb_t));
    if (!pcb) {
        sfree(ptd, PAGE_SIZE);
        return -1;
    }

    lprintf("memset");
    /* Ensure all entries are 0 and therefore not present */
    memset(ptd, 0, PAGE_SIZE);

    pcb->ptd = ptd;
    pcb->pid = pid;
    pcb->first_thread = NULL;
    pcb->prepared = 0;

    pcb_list_start = pcb;

    return 0;
}

/* TODO: To what extent should this function exist?
 * When we thread_fork, will we actually use this function? */
static int
new_tcb( int pid, int tid )
{
    pcb_t *owning_task;
    if (find_pcb(pid, &owning_task) < 0) {
        return -1;
    }

    tcb_t *tcb = malloc(sizeof(tcb_t));
    if (!tcb) {
        return -1;
    }

    owning_task->first_thread = tcb;

    /* Set tcb/pcb values */
    tcb->owning_task = owning_task;
    tcb->next_thread = NULL;
    tcb->tid = tid;

    return 0;
}

```


./kern/timer_driver.c

```

/** @file timer_driver.c
 * @brief Contains functions that implement the timer driver
 *
 * The PC timer rate is 1193182 Hz. The timer is configured to generate
 * interrupts every 10 ms. and so we round off to an interrupt every
 * 11932 clock cycles. (which is more accurate then rounding down to
 * 11931 clock cycles.
 *
 * @author Nicklaus Choo (nchoo)
 */

#include <x86/interrupt_defines.h> /* INT_CTL_PORT, INT_ACK_CURRENT */
#include <x86/asm.h> /* outb() */
#include <assert.h> /* assert() */
#include "stddef.h" /* NULL */
#include <x86/timer_defines.h> /* TIMER_SQUARE_WAVE */

#define INTERRUPT 100
#define SHORT_LSB_MASK 0x00FF
#define SHORT_MSB_MASK 0xFF00

/* Initialize tick to NULL */
static void (*application_tickback) (unsigned int) = NULL;

/* Total ticks caught */
static unsigned int total_ticks = 0;

/*****
 *
 * Internal helper functions
 *
 *****/

/** @brief Update total number of timer interrupts received and call the
 * application provided timer callback function.
 *
 * The interface for the application provided timer callback as written in the
 * handout says that 'Timer callbacks should run "quickly"', and so
 * timer_int_handler() waits for the callback application_tickback() to
 * return quickly before sending an ACK to the relevant I/O port and returning.
 *
 * @return Void.
 */
void timer_int_handler(void) {

    /* Update total ticks */
    total_ticks += 1;

    /* Pass total ticks to application callback which should run quickly */
    application_tickback(total_ticks);

    /* Acknowledge interrupt and return */
    outb(INT_CTL_PORT, INT_ACK_CURRENT);
    return;
}

/** @brief Initializes the timer driver
 *
 * TIMER_RATE is the clock cycles per second. To generate interrupts once
 * every 10 ms, we generate 1 interrupt every 10/1000 s which is 1/100 s.
 * Therefore the number of timer cycles between interrupts is:
 *
 * TIMER_RATE cycles      1 s
 * ----- x ----- = TIMER_RATE / 100 cycles
 *          s             100
 *
 * @param tickback Application provided function for callbacks triggered by

```

```

 * timer interrupts.
 * @return Void.
 */
void init_timer(void (*tickback)(unsigned int)) {

    assert(tickback != NULL);

    outb(TIMER_MODE_IO_PORT, TIMER_SQUARE_WAVE);
    short cycles_between_interrupts = (short)(TIMER_RATE / INTERRUPT);

    /* Round off */
    if ((TIMER_RATE % INTERRUPT) > (INTERRUPT / 2)) {
        cycles_between_interrupts += 1;

        assert(((cycles_between_interrupts - 1) * INTERRUPT) +
            (TIMER_RATE % INTERRUPT) == TIMER_RATE);
    } else {
        assert((cycles_between_interrupts * INTERRUPT) +
            (TIMER_RATE % INTERRUPT) == TIMER_RATE);
    }
    /* Send the least significant byte */
    short lsb = cycles_between_interrupts & SHORT_LSB_MASK;
    outb(TIMER_PERIOD_IO_PORT, lsb);

    /* Send the most significant byte */
    short msb = (cycles_between_interrupts & SHORT_MSB_MASK) >> 8;
    outb(TIMER_PERIOD_IO_PORT, msb);

    /* Set application provided tickback function */
    application_tickback = tickback;

    return;
}

```